

Anisha Reddy Chada

Ohio, USA | +1 (513) 383-1880 | chadaay0908@gmail.com | [LinkedIn](#)

PROFESSIONAL SUMMARY

Data Engineer with 5+ years of experience building scalable batch and real-time ETL/ELT pipelines, data warehouses, and Lakehouse platforms on AWS, Azure, and GCP, processing 2+ TB of data per day. Skilled in Python, SQL, Scala, Spark, Databricks, dbt, Airflow, Azure Data Factory, AWS Glue, and Kafka, with strong data modeling, performance tuning, data quality, governance, BI reporting, and CI/CD. Microsoft Azure and Databricks certified.

PROFESSIONAL EXPERIENCE

Data Engineer

May 2025 – Present

Intuit

United States

- Architected scalable ETL/ELT pipelines from source to end user with Python, PySpark, Scala, SQL, Databricks, and Delta Lake on Amazon S3, processing 2+ TB of data per day.
- Engineered batch and real-time pipelines consuming Apache Kafka streams with Spark Structured Streaming, AWS Lambda, REST and GraphQL APIs, and microservices, cutting a key pipeline from 4 hours to 75 minutes.
- Led a reusable, maintainable ELT framework with dbt, Apache Airflow, and Databricks Workflows, using AI coding assistants to speed development across 15+ production pipelines loading Amazon Redshift.
- Owned data quality, governance, and security with Unity Catalog RBAC, audit logging, lineage, schema validation, and AWS CloudWatch monitoring, achieving 98% successful pipeline completion.
- Optimized Spark, SQL, and Amazon Redshift workloads through partition pruning, caching, query tuning, and Delta Lake optimization, reducing selected analytical runtimes by 30%.
- Built CI/CD with Databricks Asset Bundles, GitHub Actions, Git and GitLab merge request reviews, Terraform, and Jira/ServiceNow ticketing workflows, automating deployments for 15+ pipelines.
- Partnered with stakeholders on requirements, architectural assessments, and sandbox POCs across AWS, Azure, and GCP, and presented insights from 2+ TB of daily data to business teams.
- Built ML and Generative AI data pipelines with Python, SQL, Amazon S3, OpenAI services, and vector embeddings, saving 10 hours of dataset preparation per development cycle.

Associate Data Engineer

Jan. 2022 – Jul. 2024

Dr. Reddy's Laboratories

India

- Developed batch and streaming ETL pipelines with Python, PySpark, and Spark Structured Streaming, ingesting pharmaceutical and manufacturing data from Azure Event Hub, IoT Hub, and SQL Server at 300+ GB per day.
- Built Azure Data Factory and Integration Runtime ingestion from 8+ on-premises and cloud sources, including Hadoop/Hive and Azure SQL, into ADLS and Databricks, saving 10 hours per reporting cycle.
- Designed conceptual, logical, and physical data models, an ODS, and data marts in Azure Synapse and SQL Server using T-SQL, stored procedures, and dbt, standardizing definitions across 10+ reports.
- Optimized large-volume Spark and Azure Synapse processing with Scala, Spark SQL, HiveQL, partitioning, and caching, reducing selected batch processing runtimes by 25%.
- Implemented incremental ETL processing and reusable transformation components across 12+ recurring pipelines, eliminating unnecessary full-data reloads.
- Orchestrated 15+ production workflows with Apache Airflow, Azure Data Factory, and Linux Bash scripts, monitoring with Azure Monitor and troubleshooting failures, dependencies, and retries.
- Created unit tests, integration tests, dbt tests, and data validation scripts, and supported User Acceptance Testing (UAT), contributing to 97% successful pipeline completion.
- Deployed changes through Azure DevOps and Jenkins CI/CD with Git pull request reviews, and defined source-to-target mappings with business analysts across 10+ integrations.

Data Analyst

Jan. 2021 – Dec. 2021

Tech Mahindra

India

- Profiled, cleaned, and rationalized structured and semi-structured client data from Azure SQL and REST APIs with Python (Pandas) and SQL in Agile teams, preparing 500+ GB of data per day.
- Automated recurring data preparation with SQL, Bash scripts, and Azure Data Factory, reducing manual data preparation by approximately 12 hours per reporting cycle.
- Built and automated Tableau and Power BI reporting on dimensional SQL datasets, presenting insights to client stakeholders across 10+ recurring reports.
- Established data quality checks covering schema validation, null checks, duplicate detection, and reconciliation across 15+ recurring workflows before production reporting.
- Supported cloud migration and modernization initiatives, translating client requirements into technical specifications and source-to-target mappings for 10+ reports.

TECHNICAL SKILLS

Languages: Python (Pandas), SQL, Scala, PySpark, Spark SQL, T-SQL, HiveQL, Bash
Big Data & Streaming: Apache Spark, Spark Structured Streaming, Databricks, Delta Lake, Hadoop, Hive, Apache Kafka, Azure Event Hub, IoT Hub
Cloud – AWS: S3, Glue, EMR, Lambda, EC2, ECS, Redshift, CloudWatch
Cloud – GCP: Google Cloud Platform (GCP)
Cloud – Azure: Data Factory, Integration Runtime, Data Lake Storage (ADLS), Synapse Analytics, Azure SQL, Azure DevOps, Azure Monitor
Data Platforms & Governance: Unity Catalog, Databricks Workflows, Databricks Asset Bundles, Lakehouse, Medallion Architecture, RBAC, Audit Logging, Data Lineage
Data Modeling & Warehousing: Conceptual, Logical, and Physical Models, Dimensional Modeling, Star Schema, Normalization, Data Marts, ODS, Data Lakes
Databases: SQL Server, Azure SQL, PostgreSQL, Amazon Redshift, Azure Synapse, MongoDB, Stored Procedures
ETL & Orchestration: ETL/ELT, dbt, Apache Airflow, Azure Data Factory, AWS Glue, Incremental Loads, REST APIs, GraphQL, Microservices
DevOps & Engineering Processes: CI/CD, GitHub Actions, GitLab, Azure DevOps, Jenkins, Git, Terraform, Infrastructure as Code, Jira, ServiceNow
Data Quality & Testing: Data Profiling, Data Cleaning, Data Validation, Unit Testing, Integration Testing, UAT, dbt Tests, Reconciliation
Performance: Query Optimization, Performance Tuning, Partition Pruning, Caching, Delta Lake Optimization
BI & AI: Power BI, Tableau, BI Automation, ML and Generative AI Data Pipelines, OpenAI Services, Vector Embeddings, AI Coding Assistants
Professional: Agile, Requirements Gathering, Stakeholder Communication, Presenting Insights, Architectural Assessments, POCs, Workstream Leadership

PROJECTS

- Cloud Lakehouse Modernization Pipeline** | *Python, PySpark, Azure Data Factory, Databricks, Delta Lake, PostgreSQL*
- Built an end-to-end ETL pipeline migrating 20+ relational tables (about 5M records) from PostgreSQL into an Azure Data Lake Lakehouse using Azure Data Factory and Databricks.
 - Implemented a medallion architecture with bronze, silver, and gold Delta Lake layers to cleanse, normalize, and conform data into a star schema with 4 fact and 6 dimension tables.
 - Added watermark-based incremental loading and source-to-target reconciliation checks, cutting daily refresh time by about 60% versus full reloads, with code versioned in GitHub.
- Multi-Source Data Integration Platform** | *Python, Apache Spark, MongoDB, PostgreSQL, Apache Airflow*
- Developed a pipeline integrating 1M+ semi-structured MongoDB documents with PostgreSQL records, flattening nested JSON into 8 normalized analytical tables with PySpark.
 - Orchestrated daily loads across 10+ Apache Airflow tasks with dependencies, retries, and failure alerts, maintaining 99% successful runs over a 30-day test period.
 - Applied duplicate detection and null checks that removed about 3% duplicate records, and designed a conceptual data model mapping entities across both sources to business processes.

EDUCATION

University of Cincinnati	Cincinnati, OH, USA
<i>Master of Science in Information Technology, GPA: 4.0</i>	<i>Dec. 2025</i>

CERTIFICATIONS

Microsoft Certified: Azure Data Engineer Associate (DP-203) – Microsoft
Databricks Certified Data Engineer Associate – Databricks